



## Weekly Research Progress Report

**Student:** Sungho Hong

**Date:** 5/26/2020

**# of hrs worked this week:** 30

Problems discussed in last week's report

### 1. Reading materials

- a. [Background of Concurrency Control](#)
- b. [Serial Safety Net Algorithm](#)

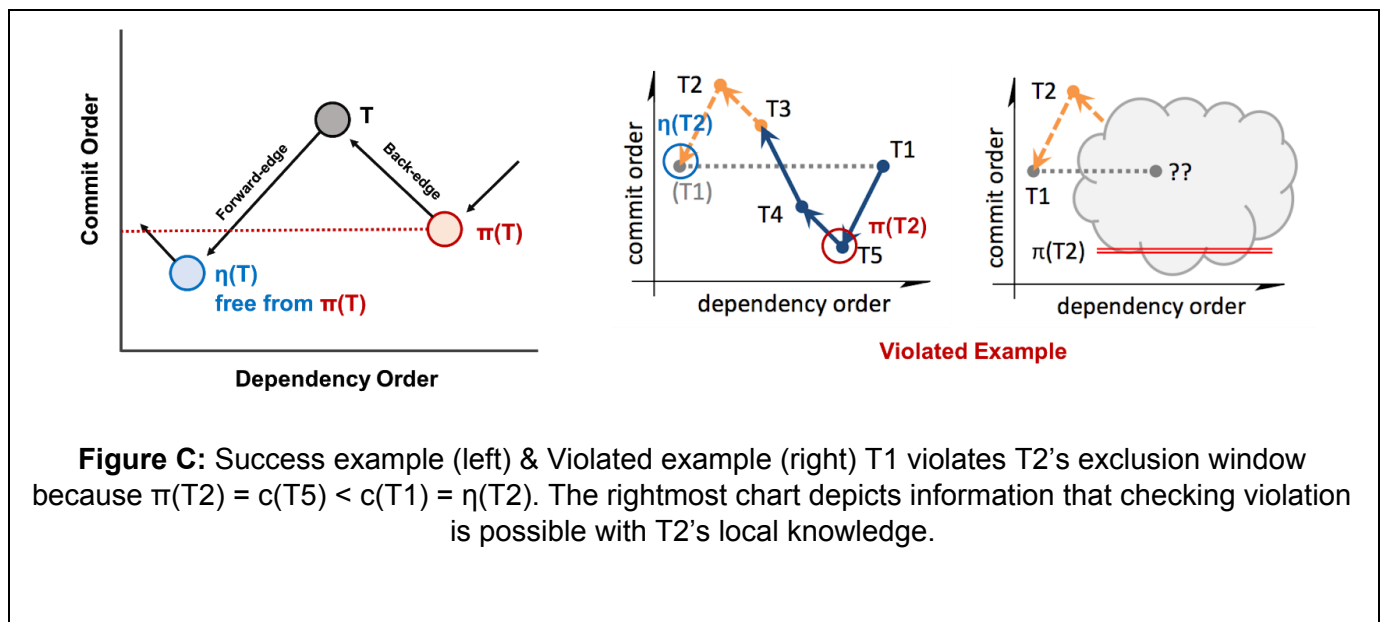
### 2. Answering questions

- a. Understand the algorithms and what it achieves
- b. Know how algorithms achieve the goal?
- c. How the algorithms can be scaled into a distributed model

## Answers

### Answering Questions

- The goal of SSN is to improve the performance of CC algorithms
  - Relaxes the strictness of CC algorithms (2PL, RC, SI) by ignoring harmless conflicts that are generally aborted by the CC algorithms
  - Enforces light-weight policy that computes from only the immediate successors of a transaction
- **Figure C:** The SSN achieves the goal by following the exclusive window policy.
  - Aborts  $\pi(T) < \eta(T)$
  - $\pi(T)$  The lowest commit time of the transaction found from the back-edges
  - $\eta(T)$  The commit time of the most recently committed transaction from the forward-edge



- Extending SSN to distributed model
  - The SSN does not assume network communication overhead
    - The SSN assumes the snapshot versions are stored in a memory of a single machine.
    - In distributed systems, additional overhead of tracking down the versions in the remote machine needs to be considered.
  - The SSN is only tested on optimistic CC mechanisms
    - The SSN requires to test the performance of both pessimistic and optimistic methods in a distributed environment.
  - The SSN does not have the propagation policy for replicated data

- For example, client-centric consistency models apply `read_set` and `write_set` to maintain consistency of the versions.
  - The client can track the missing versions when the local replica does not have the snapshot version.
  - The client can retry the request
  - Connect to another replica that may have the desired snapshot version.

### **Follow-up Questions**

1. Is the project focused on the in-memory database?
  - a. Should I focus on the concurrency control of DBMS?
  - b. Should I focus on understanding the Latch-based and compare & swap?
2. What is the SI solution that I need to follow?
  - a. There are different branches of SI, are there specific versions that you require me to focus on?

## Background of CC

- In-Memory Database
  - Relies primarily on memory for data storage
    - Pro: Eliminate the need to access disks
  - Con: Cost of transaction of acquiring a lock = Cost of accessing data
- CC ( Concurrency Control Algorithm )
  - Allows transactions to access database in multi programmed fashion
  - Preserve the illusion that each of them is executing alone on a dedicated system
- Trade off between strictness and performance of Concurrency Control algorithms
  - Strict policy (2PL, SSI)
    - Forbid many valid serializable schedules.
    - 2PL (Pessimistic) : Lock & change the current database state
    - MVCC, SI (Optimistic): First perform changes in a protected area & then change current database state
- Optimistic MVCC ( Multi Version Concurrency Control )
  - The DBMS maintain multiple physical versions of single object in the database
  - Writers don't block readers
  - Readers don't block writers
  - SI (Snapshot Isolation)
    - When a transaction starts, it sees a consistent snapshot of the database that existed when it started.
    - Write-Skew anomaly
  - SI Example: HEKATON MVCC
    - Transaction Lifecycle
      - Simulation Phase
        - Normal Processing
          - track txn read, scan, and write set
      - Validation Stage
        - Validation
          - validate reads and scans
          - If everything okay write new versions
      - Commit Stage
        - Post Processing
          - Update version timestamps
    - Transaction metadata
      - Read set
        - Physical versions that the transaction accessed
      - Write Set
        - Physical versions that the transaction created

- Scan Set
  - A set of queries to re-execute and check whether it gets the same result
- Limitations
  - Read/Scan set validation are expensive if the transactions access a lot of data
    - This depends on the workload
  - Appending new versions hurts the performance due to increased pointers
    - This part will degrade even further when the memory is located in a distributed environment.
  - Record-level conflict checks may be too coarse-grained and incur false positives
    - This part can be improved by using SSN.
- Reference
  - [Explanation of Pessimistic and Optimistic CC, Jens Dittrich](#)
  - [Multiversion concurrency control, CMU lecture series, 2020](#)

## Serial Safety Net (SSN)

### Challenges

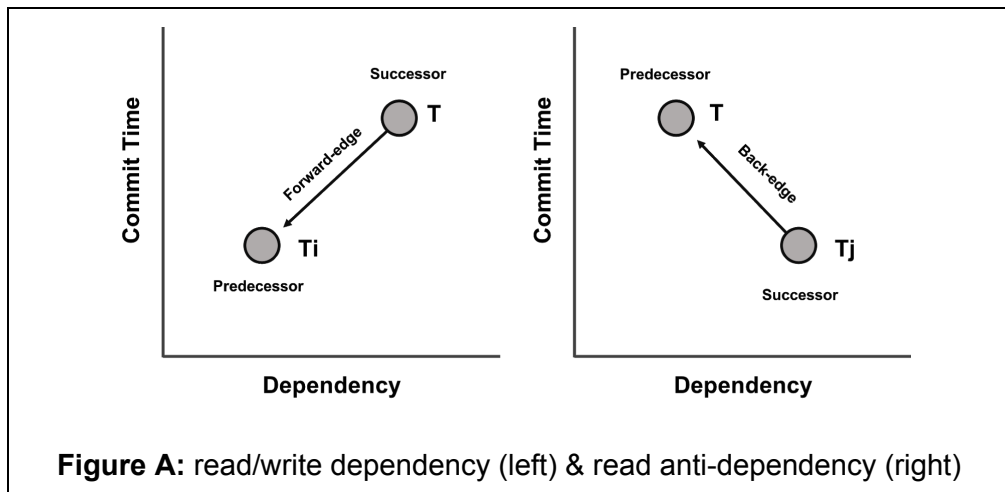
- Complexity of CC algorithms
  - bugs can lead to subtle problems
  - difficult to detect and reproduce
    - Changes of CC algorithms are not advisable
- Guarantee performance while following the CC algorithm
  - Read Committed (RC) is still the default isolation level in PostgreSQL for performance reasons
- Heterogeneous workloads
  - Read-mostly transactions.
  - Reads of stale records that are not updated recently do not have to be tracked in the transaction's read set.

### Goal of SSN

- Relax the validation stage of CC algorithm
  - CC retains control of scheduling and transactional accesses
  - SSN tracks the resulting dependencies
    - Performs a validation test by examining only direct dependencies of the committing transaction at commit time
    - Determine whether it can commit safely or must abort to avoid a potential dependency cycle

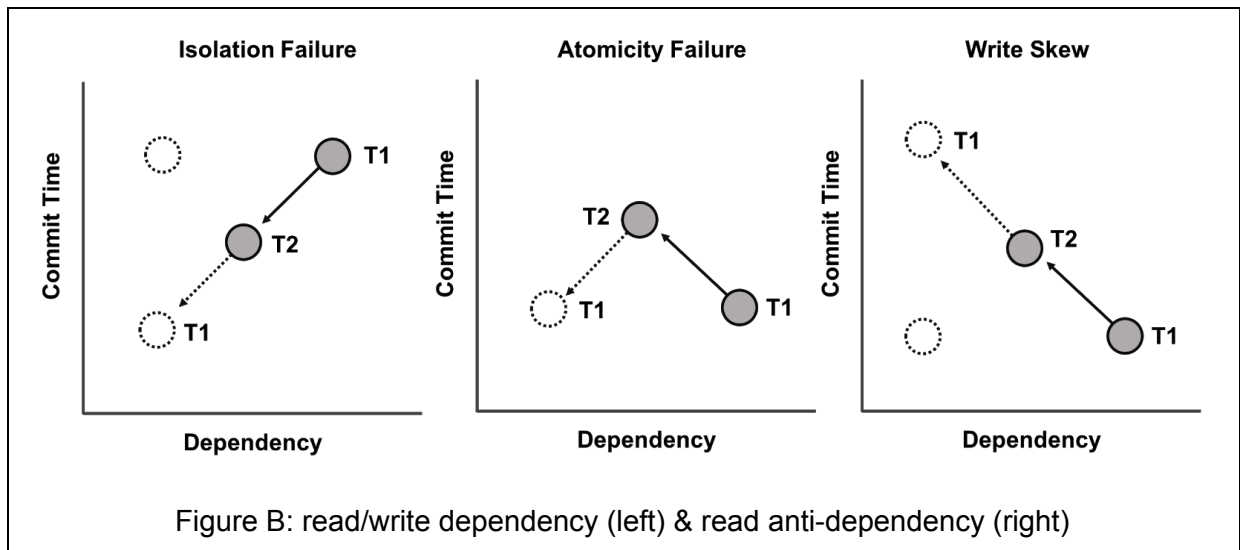
## Dependency (Figure:A)

- $T_i \xleftarrow{w:x} T$  **read/write dependency**
  - $T_i \leftarrow T$  (T depends on  $T_i$ )
    - T read or overwrote a version that  $T_i$  created
    - $T_i$  committed first and T committed next
    - $T_i$  is a predecessor of T
    - T is a successor of  $T_i$
- $T \xleftarrow{r:w} T_j$  **read anti-dependency**
  - $T \leftarrow T_j$  ( $T_j$  depends on T)
    - T read a version that  $T_j$  overwrote
    - $T_j$  committed first and T committed next
    - T is a predecessor of  $T_j$
    - $T_j$  is a successor of T



## Serialization Failures (Figure B)

- $T1 \xrightarrow{w:x} T2 \xrightarrow{w:x} T1$  **Isolation**
  - T1 and T2 saw each other's writes
- $T1 \xrightarrow{w:x} T2 \xrightarrow{r:w} T1$  **Atomicity**
  - T2 saw some but not all T1's writes
- $T1 \xrightarrow{r:w} T2 \xrightarrow{r:w} T1$  **Write Skew**
  - T1 and T2 each overwrote a value that the other read



## Isolation Levels

- Read Committed (RC)
  - Reads return the newest committed version of a record
  - Reads never block Reads
  - Writes add a new version that overwrites the latest one
  - Writes blocks only if the latter is uncommitted
- Snapshot Isolation (SI)
  - Each transaction reads from a consistent snapshot
    - newest version of each record that predates some timestamp
  - Writers must abort if they would overwrite a version created after their snapshot
- Strict Two-Phase Locking (2PL)
  - Reads return the newest version of a record, blocking if it has not committed yet
  - Writes replace the latest version, blocking if there are any in-flight reads or writes on the record by other transactions

## SSN Policies (Figure C)

- $\pi(T)$

- The lowest commit time of the transaction found from the back-edges

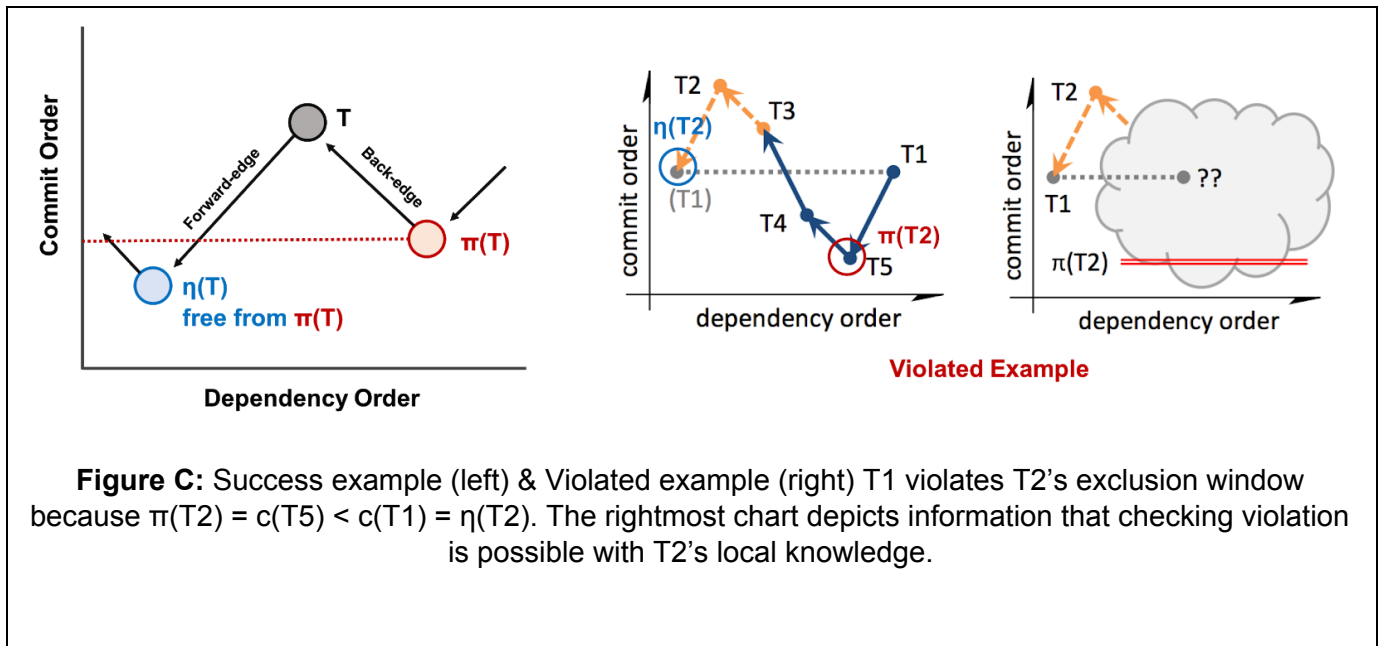
$$\begin{aligned}\pi(T) &= \min \left( c(U) : T \xleftarrow{b^*} U \right) \\ &= \min \left( \left\{ \pi(U) : T \xleftarrow{b} U \right\} \cup \{c(T)\} \right)\end{aligned}$$

- $\pi(T) < c(T)$ : The dangerous transactions committed first
- the values of  $c(T)$  and  $\pi(T)$  are fixed once  $T$  has committed
  - This would be computed from only the immediate successors of a transaction in  $G$ , without traversing the whole graph

- $\eta(T)$

- The commit time of the most recently committed transaction from the front-edge

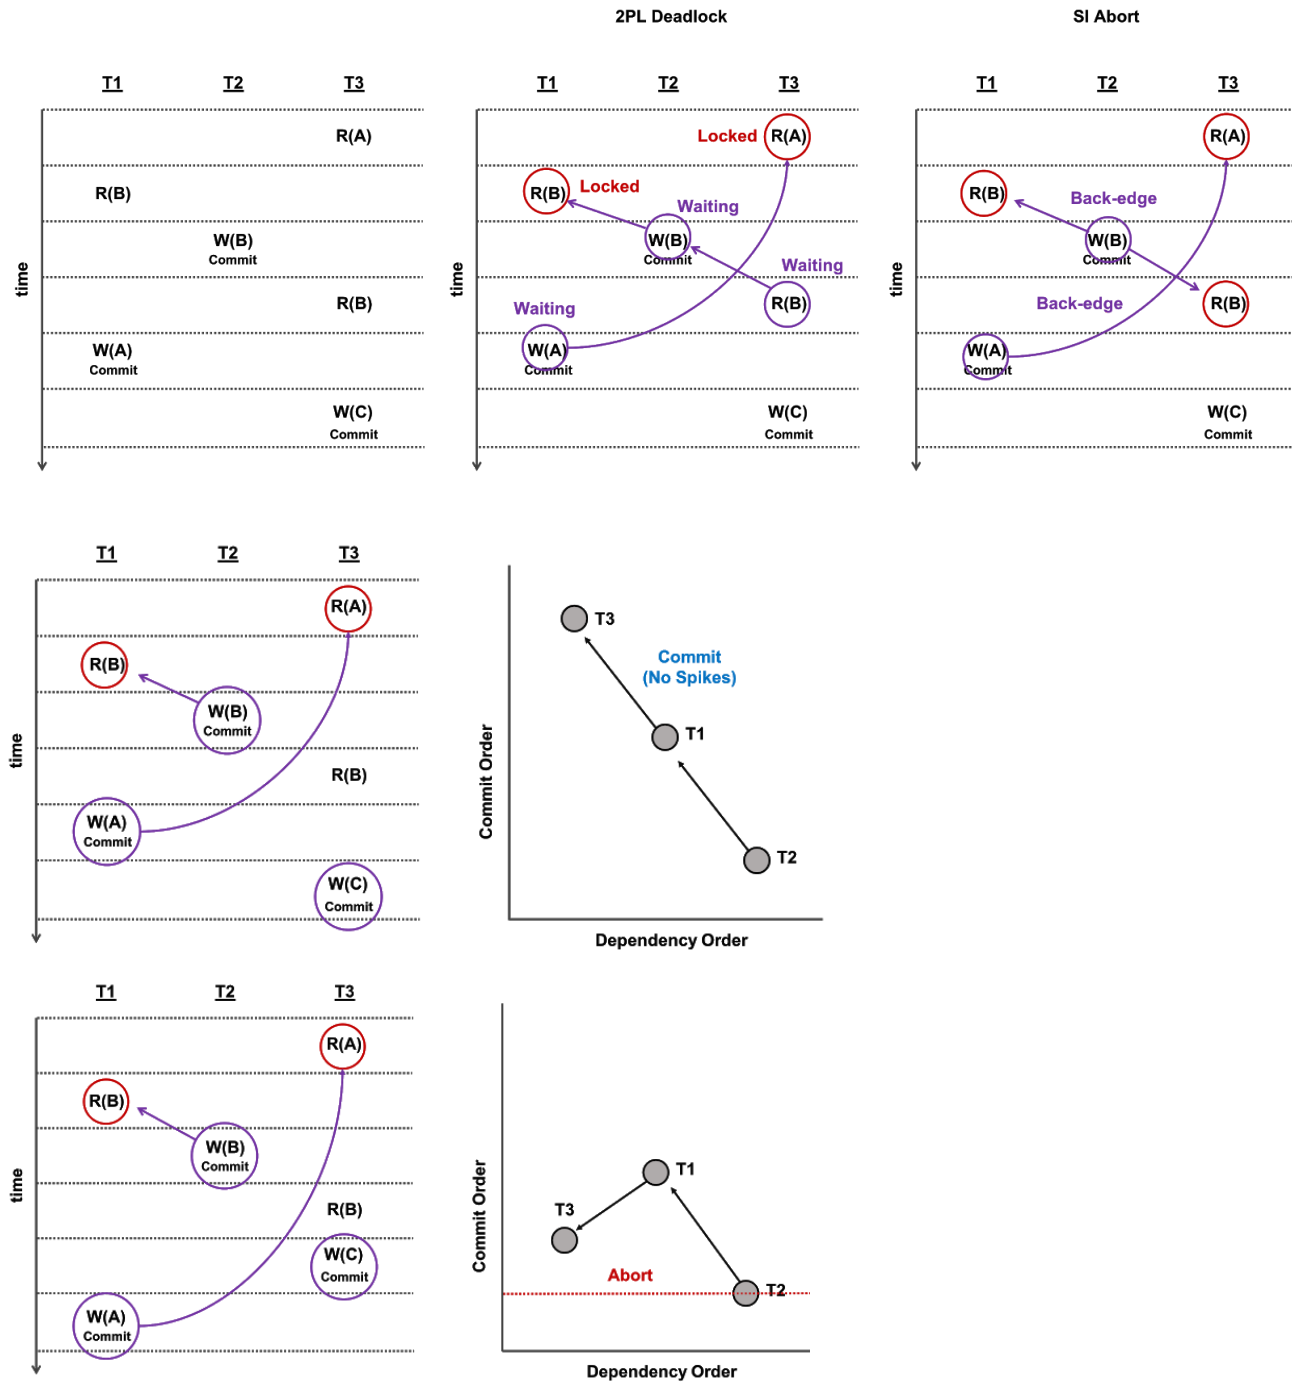
$$\eta(T) = \max \left( \left\{ c(U) : U \xleftarrow{f} T \right\} \cup \{-\infty\} \right)$$





## SSN Policies (Figure C)

- SSN ignore a large fraction of harmless back edges while still detecting all harmful ones
- Comparison among 2PL, SI, and SI + SSN



## Implementation

- Read & Write sets
  - Each transaction maintain its footprints using read and write set
  - Contain all the versions read and written by the transaction
- Worker Thread
  - Walk through the version chain to find the latest committed version that is visible to the transaction

### SSN read

```
# transaction t and snapshot v
# Reads the transaction t, and receives a reference to the appropriate version

def ssn_read(t, v):

    # if snapshot is not in the write_set
    if v not in t.writes:
        # update the  $\eta(T)$ 
        # by choosing the maximum value between  $\eta(\text{version})$  and the current timestamp
        t.pstamp = max(t.pstamp, v.cstamp)

    # if there is no pre  $\pi(T)$ 
    # if the version has not yet been overwritten,
    # it will be added to T's read set and checked for late-arriving overwrites
    # during pre-commit
    If v.sstamp is infinity
        # add the version to the read_set
        t.reads.add(v)

    # if there is  $\pi(T)$ 
    else:
        # compare the least value between the  $\pi(\text{version})$  and  $\pi(\text{transaction})$ 
        t.sstamp = min(t.sstamp, v.sstamp)

    # verifies the exclusion window and aborts if a violation is detected.
    verify_exclusion_or_abort(t)
```

### SSN write

```
# transaction t and snapshot v
def ssn_write(t, v):
    # if the snapshot is not in the write set
    if v not in t.writes:
```

```

# update the  $\eta(T)$  by choosing
# the maximum value between  $\eta(\text{version})$  and the current timestamp
# a write will never cause inbound read anti-dependencies
# a write can trigger outbound read anti-dependencies
t.pstamp = max(t.pstamp, v.prev.pstamp)

# add snapshot to the write_set
t.writes.add(v)

# remove the snapshot from the read_set
# avoid violating T's own exclusion window and trigger abort
t.reads.discard(v)

# verify the dependency cycle
verify_exclusion_or_abort(t)

```

## SSN commit

```

def ssn_commit(t):
    # T requests a commit timestamp  $c(T)$  in the in-flight status
    # T is no longer allowed to perform reads or writes
    t.cstamp = next_timestamp()

    t.sstamp = min(t.sstamp, t.cstamp)
    for v in t.reads:
        t.sstamp = min(t.sstamp, v.sstamp)

    for v in t.writes:
        t.pstamp = max(t.pstamp, v.prev.pstamp)

    # transactions having  $\eta(T) < \pi(T)$  are allowed to commit
    verify_exclusion_or_abort(t)
    t.status = COMMITTED

    # the transaction updates  $c(V)$  for each version it create
    #  $\pi(V)$  for each version it overwrote
    #  $(V)$  for each non-overwritten version it read
    for v in t.reads:
        v.pstamp = max(v.pstamp, t.cstamp)
    for v in t.writes:
        v.prev.sstamp = t.sstamp
        v.cstamp = v.pstamp = t.cstamp

```