



Weekly Research Progress Report

Student: Sungho Hong

Date:6/12/2020

of hrs worked this week: 30

Problems discussed in last week's report

1. Categorizing test cases

- a. [Test cases for RC and SI](#)
- b. [Reference materials for test-cases of RC and SI](#)

2. Reference materials

- a. Making snapshot isolation serializable, ACM Transactions on DB systems, 2005
- b. A Read-Only Transaction Anomaly Under Snapshot Isolation, SIGMOD, 2004
- c. A Critique of ANSI SQL Isolation Levels, SIGMOD, 1995

Test cases for RC and SI

O : Preventable X : Unavoidable

Test ID	P0	P1	P2	P3	P4	P5	P6
Isolation	Dirty Read	Dirty Write	Fuzzy Read	Lost update	Read Skew	Write Skew	Read Anomaly
RC	O	O	X	X	X	X	X
SI	O	O	O	O	O	X	X
RC + SSN	X	X	O	O	O	O	O

P0: Dirty Read

T1 has read a data item that was never committed and so never really existed.

One of the transactions is expected to abort in this scenario.

	5	10	15	20
T0	(X=0)	W(X=10)	C or R	
T1	R(X=10)			C or R

T0	(X=0)	W(X=10)	C or R	
T1	R(X=10)			C or R

First: T0 commits and T1 aborts as RC, **Second: T0, T1 commits**

The reason is kvstore is empty when T1 reads the data X, and the kvstore is still empty when the T0 is trying to add the data X.

P1: Dirty Write

It is unclear what the correct data value should be.

One of the transactions is expected to abort in this scenario.

	5	10	15	20
T0	W(X=10)	C or R		
T1	W(X=20)			C or R

T0	W(X=10)	C or R		
T1	W(X=20)			C or R

First, Second: Both T0 and T1 commits

P2: Fuzzy / Non-Repeatable Read

T0 receives a modified value or discovers that the data item has been deleted.

The first one is expected to abort and the second one is expected to commit.

	5	10	15	20	25
T0	R(X=10)				R(X=20) C or R
T1	W(X=20 or delete)				C or R

	5	10	15	20	25
T0	R(X=10)			R(X=10)	C or R
T1	W(X=20 or delete)				C or R

First: T0 aborts, Second: both T0 and T1 commits

P3: Lost Update

T1's update will be lost. Both of the scenarios contain anti-read dependencies, so the RC+SSN will abort automatically.

	5	10	15	20	25	30
T0	R(X=100)		W(x=130)			C
T1	R(X=100)		W(X=120)	C		

T0 aborts and T1 commits

P4: Read Skew

T1 reads y, it may see an inconsistent state. Both of the scenarios contain anti-read dependencies, so the RC+SSN will abort automatically. Fuzzy Reads is a degenerate form of Read Skew where x=y.

	5	10	15	20	25	30
T0	R(X)				R(Y)	C or A
T1	W(X)		W(Y)	C		

	5	10	15	20	25	30	35	40
T0	R(X=50)				R(y=90) C			

T1	R(x=50) W(x=10) R(y=50) W(y=90) C
----	-----------------------------------

P4: Write Skew

Two different data items were updated, each under the assumption that the other remained stable. Both of the scenarios contain anti-read dependencies, so the RC+SSN will abort automatically.

	5	10	15	20	25	30	35	40	45	50
T0	R(X)		W(Y)		C		R(X) R(Y)		C	
T1		R(Y)		W(X)		C		R(X) R(Y)		C

T0	R(X)		W(Y)		C		R(X) R(Y)		C	
T1		R(Y)		W(X)		C		R(X) R(Y)		C

	5	10	15	20	25	30	35
T0	R(X=50)	R(Y=50)			W(y=-40)		C
T1			R(X=50)	R(y=50)			C

	5	10	15	20	25	30	35	40
T0	R(X0,70)		R(Y0,80)		W(X1, -30)		C	
T1		R(X0,70)		R(Y0,80)		W(Y2, -20)		C

P5: Read-Only Anomaly

it was assumed that read-only transactions always execute serializably, without ever needing to wait or abort because of concurrent update transactions. Read only transaction T3 prints out X = 0 and Y = 20, while final values are Y = 20 and X = -11. The fact that SI allows commit order different than serial order is what causes the anomaly.

	5	10	15	20	25	30	35	40	45	50
T0			R(Y0,0)	W(Y1,20)		C				
T1	R(X0,0)	R(Y0,0)						W(X2, -11)		C
T2					R(X0,0)	R(Y1,20)		C		

Reference Materials of RC and SI(MVCC)

Isolation Levels

- trade throughput for correctness
 - Lower isolation levels increase transaction concurrency but risk showing transactions a fuzzy or incorrect database

Transaction

- a set of actions such as Reads and Writes that transform the database from one consistent state to another

History

- models the interleaved execution of a set of transactions as a linear ordering of their actions

Dependency graph

- defining the temporal data flow among transactions.
- Two histories are equivalent if they have the same committed transactions and the same dependency graph

Serializability

- A history is serializable if it is equivalent to a serial history
 - if the history has the same dependency graph (inter-transaction temporal data flow) as some history that executes transactions one at a time in sequence

Concurrency

- if T1 and T2 transactional lifetimes overlap
 - $[\text{start}(T1), \text{commit}(T1)] \cap [\text{start}(T2), \text{commit}(T2)] \neq \emptyset$.
 - writes by concurrent transactions, are not visible to the transaction
 - When T_i is ready to commit, it obeys the First Committer Wins rule

First Committer Wins

- T_i will successfully commit if and only if no concurrent transaction T_k has already committed writes (updates) of rows or index entries that T_i intends to write.