



Weekly Research Progress Report

Student: Sungho Hong

Date: 7/13/2020

of hrs worked this week: 30

Problems discussed in last week's report

1. [Exploring Jepsen I/O](#)

- a. Example of testing etcd with Jepsen I/O
- b. Possibility of integrating Jepsen to RamCloud
- c. Checking anomalies with Jepsen

2. Debugging anomalies from ClusterPerf

- a. Debugged results

Elle: Inferring Isolation Anomalies from Experimental Observations

Jepsen 2020

Background

- Existing solutions require hand-proven invariants.
- General checkers verify linearizability
- Serializability checking is NP-complete
 - cannot use real-time constraints to reduce the search space

Challenge

- Database may not have any concept of a version order
- Database may not expose that ordering information to clients
- Cannot deduce the version in real-time order
 - Optimistic and multi-version may commit earlier versions later in time.
- Cannot identify write-write dependency
 - Database could have executed T1 and T2 in either order with equivalent

Solution

- **Elle == Isolation checker**
 - Introduce **recoverability and traceability** to infer dependencies between observed transactions
 - Elle infers a transaction dependency graph over experimentally recorded histories, and searches for cycles in that graph.
 - G0 (dirty write),
 - G1a (aborted read), G1b (intermediate read)

- G1c (cyclic information flow)
- G-single (read skew)
- G2-item (anti-dependency cycle)

- **Counter & Sets**

$T_0: read(x, \{0\})$

$T_1: add(x, 1)$

$T_2: add(x, 2)$

$T_3: read(x, \{0, 1, 2\})$

- Read-read dependency: $T_0 \rightarrow T_3$
- Write-read dependency: $T_1 \rightarrow T_3$, $T_2 \rightarrow T_3$
- Anti-read dependency: $T_0 \rightarrow T_1$ $T_0 \rightarrow T_2$

- **List**

- add order to our values by letting each version be a list
- a write appends a unique value
- Write-write dependency:
 - If there is a read of $\{0, 1\}$ or $\{0, 2\}$

- **Traceability**


- Restricting our histories to the objects with list
- Get the largest version of the object
- Reason about the version order using the results of individual read operations

- **Recoverability**

- Identify aborted and intermediate versions
- Restricting observations to a single write per object makes recovery trivial.
 - Counters and sets are difficult to recover in general: a set like $\{1, 2\}$ could have resulted either from a write of 1 or 2.

Background

- PostgreSQL added support for serializable snapshot isolation (SSI)
- SSI extends SI by checking for dangerous structures
 - a pair of adjacent read-write dependencies between three transactions.



- Preventing these dangerous structures, in addition to snapshot isolation's normal rules, yields only serializable executions

Challenge



- Transactions executed with serializable isolation on a single PostgreSQL instance were not serializable

Design

- Kill PostgreSQL processes in random order
- Generate randomized transactions of append and read operations
- Use Elle isolation checker to infer a transaction dependency graph over recorded histories

Results

- PostgreSQL allows write-skew, and dangerous structures



Three transactions involve a read-only transaction which precedes, via two adjacent rw anti-dependencies, a transaction which wrote state which the read-only transaction observed.

Conclusion

- This code has gone essentially untouched since the introduction of serializable snapshot isolation in 2011

Integration with RamCloud

1. Using Jepsen + Elle
 - a. Implement the client for Ramcloud
 - b. Additionally benefit from intensive test-case scenarios from Jepsen
2. Using Elle (bypass Jepsen)
 - a. Record the history and convert them into Jepsen history
 - b. Can use the test cases provided by RamCloud

Reference Materials

Comparison between Linearizability and Serializability

| Linearizability | Serializability |
|---|--|
| <ul style="list-style-type: none">- Based on time- Deterministic- Provides a real-time (i.e., wall-clock) guarantee on the behavior of a set of single operations | <ul style="list-style-type: none">- Program order- Freedom to interleave operations as long as the ordering from each client is preserved- Guarantees that the execution of a set of transactions is equivalent to some serial execution (total ordering) of the transactions. |